

R-Trees

A Dynamic Index Structure for Spatial Searching

Based on Antonin Guttman's 1984 Paper · *February 23, 2026*



PRESENTED BY

Farhan Labib

2205071

Fatima Sad Sudipta

2205063

Bangladesh University of
Engineering & Technology

Outline

- Introduction
- R-Tree Structure
- Search Simulation
- Insertion & Node Splitting
- Full Insertion Walkthrough
- Deletion
- Key Takeaways



Introduction

The Problem

Traditional Indexes

- B-trees, hash tables
- One-dimensional keys
- Exact-match queries

They struggle with...

- Multi-dimensional data
- Objects with *spatial extent*
- “Find all objects overlapping S ”

Real-world examples

- GIS: counties, roads, buildings
- CAD: circuit layouts
- Games: collision detection
- Maps: nearest-place search

We need a better structure!

The Problem

Traditional Indexes

- B-trees, hash tables
- One-dimensional keys
- Exact-match queries

They struggle with...

- Multi-dimensional data
- Objects with *spatial extent*
- “Find all objects overlapping S ”

Real-world examples

- GIS: counties, roads, buildings
- CAD: circuit layouts
- Games: collision detection
- Maps: nearest-place search

We need a better structure!

The Problem

Traditional Indexes

- B-trees, hash tables
- One-dimensional keys
- Exact-match queries

They struggle with...

- Multi-dimensional data
- Objects with *spatial extent*
- “Find all objects overlapping S ”


Real-world examples

- GIS: counties, roads, buildings
- CAD: circuit layouts
- Games: collision detection
- Maps: nearest-place search

We need a better structure!

R-Tree Structure

Core Idea: Minimum Bounding Rectangles




Step 1 — Objects

Four spatial objects in 2-D space.

- Every tree node \approx one disk page
- $m \leq$ entries $\leq M$ per node
- All leaves at the same depth

Core Idea: Minimum Bounding Rectangles




Step 2 — Group left pair

Draw the tightest rectangle around A and B: the **Minimum Bounding Rectangle (MBR)**.

- Every tree node \approx one disk page
- $m \leq$ entries $\leq M$ per node
- All leaves at the same depth

Core Idea: Minimum Bounding Rectangles



Step 3 — Group right pair

Each MBR becomes one **entry** in an R-tree internal node. MBRs *may* overlap.

- Every tree node \approx one disk page
- $m \leq$ entries $\leq M$ per node
- All leaves at the same depth


R-Tree Node Types & Tree Shape

Leaf node entry (I , $tuple\text{-}id$)

- I = MBR tightly wrapping the object
- $tuple\text{-}id$ = pointer to data record


Internal node entry (I , $child\text{-}ptr$)

- I = MBR covering *all* children's MBRs
- $child\text{-}ptr$ = address of child page



Search Simulation

Simulation: Search — Step 1, Tree Structure




Step 1 — Tree structure

Three leaf groups: **R8**, **R9**, **R10**.
Each group's MBR tightly wraps its objects.

[

ghost]Step 2 — Coming next Issue
a spatial query box $S \dots$

Simulation: Search — Step 2, Issue Query S




Step 2 — Issue query S

For each MBR, ask:
Does it overlap S ?

Step 3 — Coming next

R8 and R9 will match...

Simulation: Search — Step 3, R8 & R9 Match



Step 3 — R8 and R9 match


Both MBRs overlap *S*.

Descend into both subtrees and report overlapping objects.

Step 4 — Coming next

R10 will be pruned entirely...

Simulation: Search — Step 4, R10 Pruned



Step 4 — R10 pruned

$$R10 \cap S = \emptyset$$

Skip the entire R10 subtree.
Zero extra disk reads!

Key lesson

Good MBRs \Rightarrow more pruning
 \Rightarrow fewer disk reads
 \Rightarrow faster queries

Insertion & Node Splitting

Insertion — Overview

1. **ChooseLeaf** — pick child needing **least** MBR enlargement
2. Insert entry into chosen leaf
3. If leaf has $> M$ entries: **SplitNode**
4. **AdjustTree** — update ancestor MBRs upward
5. If root split: create **new root** one level higher

The critical step

Quality of **node splitting** drives future search performance.

Goal: small, non-overlapping MBRs.

Insertion — Overview

1. **ChooseLeaf** — pick child needing **least** MBR enlargement
2. **Insert** entry into chosen leaf
3. If leaf has $> M$ entries: **SplitNode**
4. **AdjustTree** — update ancestor MBRs upward
5. If root split: create **new root** one level higher

The critical step

Quality of **node splitting** drives future search performance.

Goal: small, non-overlapping MBRs.

Insertion — Overview

1. **ChooseLeaf** — pick child needing **least** MBR enlargement
2. **Insert** entry into chosen leaf
3. If leaf has $> M$ entries: **SplitNode**
4. **AdjustTree** — update ancestor MBRs upward
5. If root split: create **new root** one level higher

The critical step

Quality of **node splitting** drives future search performance.
Goal: small, non-overlapping MBRs.

Insertion — Overview

1. **ChooseLeaf** — pick child needing **least** MBR enlargement
2. **Insert** entry into chosen leaf
3. If leaf has $> M$ entries: **SplitNode**
4. **AdjustTree** — update ancestor MBRs upward
5. If root split: create **new root** one level higher

The critical step

Quality of **node splitting** drives future search performance.

Goal: small, non-overlapping MBRs.


Insertion — Overview

1. **ChooseLeaf** — pick child needing **least** MBR enlargement
2. **Insert** entry into chosen leaf
3. If leaf has $> M$ entries: **SplitNode**
4. **AdjustTree** — update ancestor MBRs upward
5. If root split: create **new root** one level higher

The critical step

Quality of **node splitting** drives future search performance.
Goal: small, non-overlapping MBRs.

Simulation: ChooseLeaf — Step 1, New Object Arrives




Step 1 — Obj7 arrives

Must find the best leaf node to insert into.

Step 2 — Coming next

Compute enlargements for each group...

Simulation: ChooseLeaf — Step 2, Compare Enlargements



Step 2 — Enlargements

R8: large growth needed


R9: minimal growth ✓

R10: large growth needed

Step 3 — Coming next

Insert into R9 and update MBR...

Simulation: ChooseLeaf — Step 3, Insert into R9



Step 3 — Inserted!

$R9 = \{Obj3, Obj4, Obj7\}$

3 entries = M . No split.

MBR updated to wrap Obj7.

Node Splitting — Three Algorithms

Exhaustive

All 2^M groupings.
Optimal quality.
 $O(2^M)$ — too slow.

Quadratic (*focus*)

Heuristic seeds.
Good quality.
 $O(M^2)$ per split.


Linear

Axis extremes.
Acceptable quality.
 $O(M)$ — fastest.

Goal of any split

Two nodes with **small, minimally overlapping** MBRs.
Better split \Rightarrow better search pruning \Rightarrow fewer disk reads.

Quadratic Split — Step 1: PickSeeds, Meet the Entries



Overflowing node

4 entries: A, B, C, D ($> M = 3$).

Must split into two groups.

PickSeeds formula

For every pair (E_i, E_j) :

$$d = \text{area}(E_i \cup E_j) - \text{area}(E_i) - \text{area}(E_j)$$

Choose pair with **largest d** (most wasted space).

Quadratic Split — Step 2: PickSeeds, Worst Pair Found



Seeds: A and D

Diagonally opposite — their joint MBR wastes the most space.


$$s_1 = A \rightarrow G_1$$

$$s_2 = D \rightarrow G_2$$

Next — PickNext

Assign B and C to their preferred group...

Quadratic Split — Step 3: PickNext, Assign Remaining



Rule

$$\delta_i = \Delta\text{area}(G_1) - \Delta\text{area}(G_2)$$

Assign entry with **largest $|\delta_i|$** first.


Round 1 result

B strongly prefers G_1 (near A).

C strongly prefers G_2 (near D).

Both assigned. **Split complete!**

Quadratic Split — Step 4: Final Groups



Result


Node 1: {A, B} — left cluster
Node 2: {C, D} — right cluster

Compact, non-overlapping MBRs.

Adjust Tree

Parent gets two new entries pointing to Node 1 and Node 2.
Ancestor MBRs updated upward.


Good Split vs Bad Split



Bad Split

Large overlap

⇒ query visits BOTH branches



Good Split


No overlap

⇒ query prunes one branch entirely

Full Insertion Walkthrough

Walkthrough — Setup


R-tree with $M = 3$, starting state:



Plan: insert **Obj7** (near R9) — no overflow.

Then insert **Obj8** (near R9) — triggers **overflow + split**.

Walkthrough — Insert Obj7, R9 Updated




ChooseLeaf → R9

R9 needs minimum enlargement.


Result — No Split

$R9 = \{Obj3, Obj4, Obj7\}$

3 entries = M . Tree unchanged.



Walkthrough — Insert Obj8, Overflow!



Overflow

R9 would have 4 entries $> M = 3$.
SplitNode triggered.


Quadratic split result

Seeds: Obj3 & Obj7 (max wasted area).
R9a = {Obj3, Obj8}
R9b = {Obj4, Obj7}

AdjustTree

Root gains R9a and R9b.
4 children total — still $\leq M$.

Walkthrough — Final Tree After Split



Root now has 4 children. Tree height unchanged. All MBRs compact thanks to the quadratic split.

Deletion

Deletion Algorithm

1. **FindLeaf** — locate the leaf containing E
2. Remove E from that leaf
3. **CondenseTree** — if node has $< m$ entries:
 delete it; **re-insert** orphaned entries
4. Propagate MBR shrinkage upward
5. If root has one child: shrink tree height

Why re-insert, not merge?

- Reuses Insert logic
- Entries find natural home
- Gradually **improves** structure
- Simpler than spatial merge

Deletion Algorithm

1. **FindLeaf** — locate the leaf containing E
2. **Remove E** from that leaf
3. **CondenseTree** — if node has $< m$ entries:
 delete it; **re-insert** orphaned entries
4. Propagate MBR shrinkage upward
5. If root has one child: shrink tree height

Why re-insert, not merge?

- Reuses Insert logic
- Entries find natural home
- Gradually **improves** structure
- Simpler than spatial merge

Deletion Algorithm

1. **FindLeaf** — locate the leaf containing E
2. **Remove E** from that leaf
3. **CondenseTree** — if node has $< m$ entries:
 delete it; **re-insert** orphaned entries
4. Propagate MBR shrinkage upward
5. If root has one child: shrink tree height

Why re-insert, not merge?

- Reuses Insert logic
- Entries find natural home
- Gradually **improves** structure
- Simpler than spatial merge

Deletion Algorithm

1. **FindLeaf** — locate the leaf containing E
2. **Remove E** from that leaf
3. **CondenseTree** — if node has $< m$ entries:
 delete it; **re-insert** orphaned entries
4. Propagate MBR shrinkage upward
5. If root has one child: shrink tree height

Why re-insert, not merge?

- Reuses Insert logic
- Entries find natural home
- Gradually **improves** structure
- Simpler than spatial merge

Deletion Algorithm

1. **FindLeaf** — locate the leaf containing E
2. **Remove E** from that leaf
3. **CondenseTree** — if node has $< m$ entries:
 delete it; **re-insert** orphaned entries
4. Propagate MBR shrinkage upward
5. If root has one child: shrink tree height

Why re-insert, not merge?

- Reuses Insert logic
- Entries find natural home
- Gradually **improves** structure
- Simpler than spatial merge

Key Takeaways

Summary

What we covered

1. MBRs group nearby spatial objects
2. Search prunes non-overlapping subtrees
3. ChooseLeaf: least-enlargement heuristic
4. Quadratic split: PickSeeds + PickNext
5. Split quality drives search efficiency
6. Deletion re-inserts orphaned entries

The golden rule

Minimise **overlap** and **area** of MBRs at every split.

Used in practice

PostgreSQL/PostGIS, QGIS, ArcGIS, game engines, GEOS, Shapely, SQLite R*

References

-  Antonin Guttman
R-Trees: A Dynamic Index Structure for Spatial Searching
ACM SIGMOD, 1984
-  Wikipedia contributors
R-tree — en.wikipedia.org/wiki/R-tree



Thank You

Questions & Discussion

